

Integrating Machine Learning and Optimization for Inventory Management

Ziang Liu

Faculty of Environmental, Life, Natural Science and Technology
Okayama University

MOAI, January 26, 2024

Table of Contents

- 1 Reinforcement Learning for Inventory Optimization
 - Newsvendor Problem and Multi-armed Bandit Problem
 - Multi-Period Inventory Problem and Markov Decision Process
 - Multi-Period Inventory Problem and TD Learning
 - Inventory Management and Deep Reinforcement Learning
- 2 Surrogate Models in Optimization
 - Surrogate-based Optimization
 - Trust Region Methods
 - Bayesian Optimization
 - Surrogate-assisted Evolutionary Optimization
- 3 Appendix

References

Books

- Sutton, Richard S., and Andrew G. Barto. 2018. Reinforcement Learning, Second Edition: An Introduction. MIT Press.
- Snyder, Lawrence V., and Zuo-Jun Max Shen. 2019. Fundamentals of Supply Chain Theory. John Wiley & Sons.
- Puterman, Martin L. 1990. "Chapter 8 Markov Decision Processes." In Handbooks in Operations Research and Management Science, 2:331-434. Elsevier.
- Kochenderfer, Mykel J., and Tim A. Wheeler. 2019. Algorithms for Optimization. MIT Press.
- Jin, Yaochu, Handing Wang, and Chaoli Sun. 2021. Data-Driven Evolutionary Optimization: Integrating Evolutionary Computation, Machine Learning and Data Science. Springer Nature.







Webpages

- "UCL Course on RL." 2019. David Silver. December 19, 2019. <https://www.davidsilver.uk/teaching/>.
- Snyder, Larry. n.d. RL for Inventory Optimization. Github. Accessed January 4, 2024. <https://github.com/LarrySnyder/RLforInventory>.
- Poupart, Pascal. n.d. "CS885 Spring 2020 - Reinforcement Learning." Accessed January 4, 2024. <https://cs.uwaterloo.ca/ppoupart/teaching/cs885-spring20/schedule.html>.

Source Code

Source code is available on GitHub.

<https://github.com/zi-ang-liu/ML-and-Opt-for-Inventory>



Mindmap

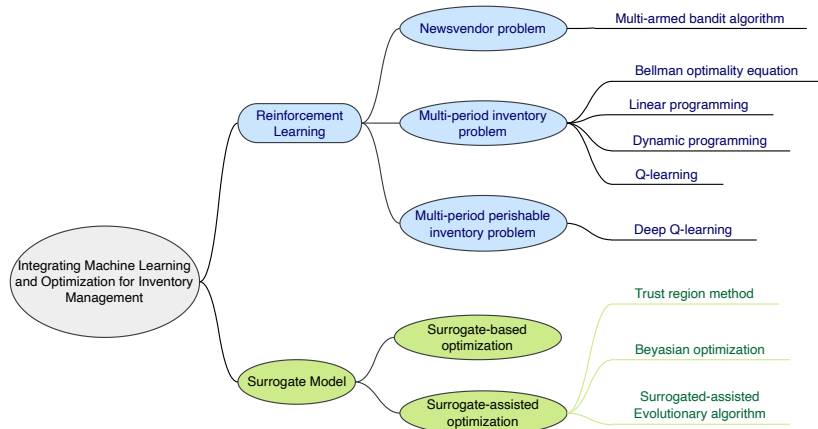


Table of Contents

- 1 Reinforcement Learning for Inventory Optimization
 - Newsvendor Problem and Multi-armed Bandit Problem
 - Multi-Period Inventory Problem and Markov Decision Process
 - Multi-Period Inventory Problem and TD Learning
 - Inventory Management and Deep Reinforcement Learning
- 2 Surrogate Models in Optimization
 - Surrogate-based Optimization
 - Trust Region Methods
 - Bayesian Optimization
 - Surrogate-assisted Evolutionary Optimization
- 3 Appendix

The newsvendor problem: Description

- A newspaper vendor must decide how many copies of a newspaper to order each morning
- The demand is uncertain
- Overage cost occurs when the demand is less than the order quantity
- Underage cost occurs when the demand is greater than the order quantity
- What is the optimal order quantity?

The newsvendor problem: Formulation

Parameters:

- D : Demand (random variable)
- h : Holding cost (overage cost)
- p : Stockout cost (shortage cost)

Decision variables:

- Q : Order quantity

Objective:

- $g(Q) = \mathbb{E} [h(Q - D)^+ + p(D - Q)^+]$
- $(x)^+ = \max\{x, 0\}$

The newsvendor problem: Optimization

Objective function:

$$\begin{aligned} & \mathbb{E} [h \max \{Q - D, 0\} + p \max \{D - Q, 0\}] \\ &= h \int_0^Q (Q - d)f(d)dd + p \int_Q^\infty (d - Q)f(d)dd \end{aligned}$$

First order derivative: $hF(Q) - p(1 - F(Q))$

Second order derivative: $hf(Q) + pf(Q)$

Optimal order quantity: $Q^* = F^{-1} \left(\frac{h}{h+p} \right)$

The newsvendor problem: Example

■ $D \sim \mathcal{N}(50, 8^2)$

■ $h = 0.18, p = 0.7$




Figure: Continuous newsvendor problem

The discrete newsvendor problem: Example




Figure: Discrete newsvendor problem

$g(Q)$ is still convex.

Multi-armed bandit problem

Problem statement:

- k different actions we can take
- After each action, we receive a reward from a stationary probability distribution
- The reward distribution is unknown
- Objective: maximize the total reward over some time period

The k-armed bandit problem: Example

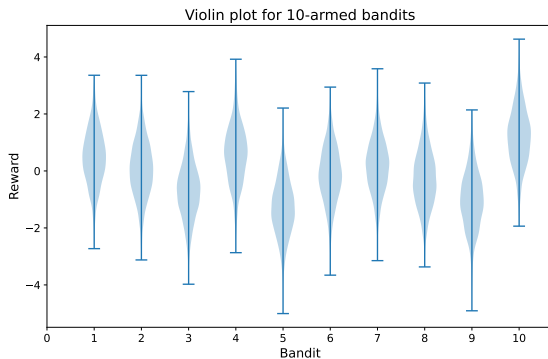


Figure: The 10-armed bandit problem

The k-armed bandit problem: Notations

- A_t : Action at time t
- R_t : Reward at time t
- $q_*(a)$: Expected reward for taking action a ,
 $q_*(a) = \mathbb{E}[R_t | A_t = a]$
- $Q_t(a)$: Estimated value of action a at time t ,
$$Q_t(a) = \frac{\sum_{i=1}^{t-1} R_i \mathbb{I}\{A_i = a\}}{\sum_{i=1}^{t-1} \mathbb{I}\{A_i = a\}}$$

$\mathbb{I}\{A_i = a\}$ is an indicator function, $\mathbb{I}\{A_i = a\} = 1$ if $A_i = a$, otherwise $\mathbb{I}\{A_i = a\} = 0$

Incremental implementation of sample averages

$$\begin{aligned}Q_{n+1} &= \frac{R_1 + R_2 + \cdots + R_n}{n} \\&= \frac{1}{n} \sum_{i=1}^n R_i \\&= \frac{1}{n} \left(R_n + \sum_{i=1}^{n-1} R_i \right) \\&= \frac{1}{n} \left(R_n + (n-1) \frac{R_1 + R_2 + \cdots + R_{n-1}}{n-1} \right) \\&= \frac{1}{n} (R_n + (n-1)Q_n) \\&= Q_n + \frac{1}{n} (R_n - Q_n)\end{aligned}$$

The k-armed bandit problem: ϵ -greedy policy

- with probability $1 - \epsilon$, choose the action with the highest estimated value, $A_t = \arg \max_a Q_t(a)$
- with probability ϵ , choose an action randomly

A simple bandit algorithm

Algorithm 1: A simple bandit algorithm

```
1 Initialize, for  $a = 1, \dots, k$ :  
2    $Q(a) \leftarrow 0$   
3    $N(a) \leftarrow 0$   
4 for  $t = 1, 2, \dots$  do  
5   Choose  $A_t$  using  $\epsilon$ -greedy policy based on  $Q_t$   
6    $R_t \leftarrow \text{bandit}(A_t)$   
7    $N(A_t) \leftarrow N(A_t) + 1$   
8    $Q(A_t) \leftarrow Q(A_t) + \frac{1}{N(A_t)} (R_t - Q(A_t))$   
9 end
```

Solving newsvendor problem as a k-armed bandit problem

discrete newsvendor problem

- $h = 0.18$, $p = 0.7$, $D \sim \mathcal{N}(5, 1^2)$, D is discrete
- Optimal order quantity: 6, Expected Cost: 0.24446

simple bandit algorithm ($\epsilon = 0.01$, $k = 10$, $T = 2000$)

Optimal Action	Q Value
6	-0.24113
6	-0.24389
6	-0.24330
6	-0.24549
6	-0.23680

Table: Optimal actions and corresponding Q values

k-armed bandit problem and newsvendor problem

Summary

- k-armed bandit problem is nonassociative
- Only need to find the best action
- The newsvendor problem: Perishable, Lost sales
- The newsvendor problem can be solved as k-armed bandit problem

Next: Associative problem

- need to learn a policy (a mapping from states to actions)
- Multi-period inventory problem

Inventory Policy

- **Constant order quantity policy:** Order Q units at each period.
- **Base-stock policy:** Order up to S units when the inventory level is less than S , otherwise do not order.
- **(s, S) policy:** Order up to S units when the inventory level is less than s , otherwise do not order.

Inventory policy is a mapping from inventory level to order quantity.

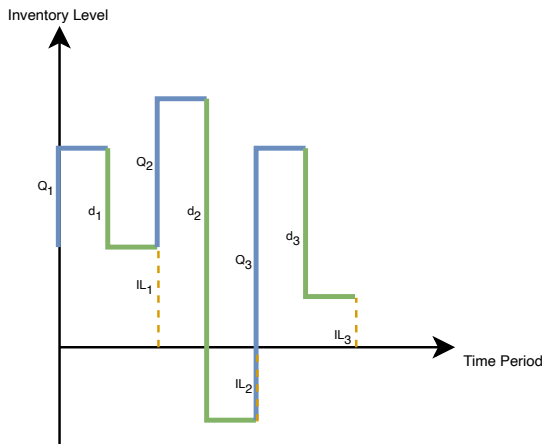
The multi-period inventory problem: Description

- Inventory is non-perishable
- Decide how much to order at each morning
- The demand is uncertain
- Unsatisfied demand is backordered
- What is the optimal order **policy**?

Notations

- D_t : Demand at period t
- h : Holding cost (overage cost)
- p : Stockout cost (shortage cost)
- Q_t : Order quantity at period t
- IL_t : Inventory level at the end of period t
 - $IL_t > 0$: Inventory on hand
 - $IL_t < 0$: Backorder
 - $IL_t = IL_{t-1} + Q_t - D_t$

The multi-period inventory problem



Objective

Cost at period t

- $g_t(Q_t) = h(IL_{t-1} + Q_t - d_t)^+ + p(d_t - IL_{t-1} - Q_t)^+$

Objective

- Total cost: $\sum_{t=1}^T g_t(Q_t)$
- Total cost, infinite horizon: $\sum_{t=1}^{\infty} g_t(Q_t)$
- Cumulative discounted cost, infinite horizon: $\sum_{t=1}^{\infty} \gamma^{t-1} g_t(Q_t)$

Problem statement

- Objective: minimize the discounted cumulative cost,
$$\sum_{t=1}^{\infty} \gamma^{t-1} g_t(Q_t)$$
- Optimal inventory decision at period t depends on the ending inventory level in the previous period, I_{t-1}
- We want to learn a policy π that maps I_{t-1} to Q_t at each period t .
- The multi-period inventory problem is an associative problem

Markov Decision Process (MDP)

Definition

A Markov Decision Process (MDP) is a tuple (S, A, P, R, γ)

- S : a set of states
- A : a set of actions
- P : transition probability,
 $p(s', r|s, a) = \mathbb{P}[S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a],$
 $p : S \times R \times S \times A \rightarrow [0, 1]$
- R : a reward function, $r(s, a) = \mathbb{E}[R_t | S_{t-1} = s, A_{t-1} = a],$
 $r : S \times A \rightarrow \mathbb{R}$
- γ : a discount factor, $\gamma \in [0, 1]$

Multi-period inventory problem as an MDP

Multi-period inventory problem as an MDP

- $S = \{IL_t \in \mathbb{Z} : IL_t \text{ is a ending inventory levels}\}$
- $A = \{Q_t \in \mathbb{N} : Q_t \text{ is a order quantity}\}$
- $p(s', r | s, a) = \mathbb{P}[D_t = s + a - s']$
- R : a reward function,
 $r(s, a) = -\mathbb{E}[h(s + a - D_t)^+ + p(D_t - s - a)^+]$
- γ : a discount factor, $\gamma \in [0, 1]$

Multi-period inventory problem as an MDP

Returns

- discount return: $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$

Policies and value functions

- policy: $\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$
- state value function: $v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s]$
- action value function: $q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$

Optimal policies and value functions

Policies

- $\pi \geq \pi'$ if $v_\pi(s) \geq v_{\pi'}(s), \forall s \in S$
- At least one optimal policy exists, denoted as π_*

Value functions

- optimal state value function: $v_*(s) = \max_\pi v_\pi(s), \forall s \in S$
- optimal action value function: $q_*(s, a) = \max_\pi q_\pi(s, a), \forall s \in S, a \in A$

Computing optimal policies and value functions

- Bellman optimality equation
- Linear programming
- Dynamic programming
 - Policy iteration
 - Value iteration


Bellman optimality equation (Optional)

$$\begin{aligned} v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\ &= \max_a \mathbb{E}_{\pi_*} [G_t | S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a] \\ &= \max_a \mathbb{E}_{\pi_*} [R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a] \\ &= \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')] \end{aligned}$$

Solving MDP: Bellman optimality equation

Bellman optimality equation

$$v_*(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_*(s')]$$



Solving MDP: Linear programming

- A less frequently used method for solving MDP
- Idea:
 - If $v(s) \geq r(s, a) + \gamma \sum_{s'} p(s'|s, a)v(s')$ for all $s \in S$ and $a \in A$, then $v(s)$ is an upper bound on $v_*(s)$
 - $v_*(s)$ must be the smallest such solution

Solving MDP: Linear programming

Linear programming formulation

$$\begin{aligned} & \text{minimize } \sum_{s \in S} \alpha_s v(s) \\ & \text{s.t. } v(s) \geq r(s, a) + \gamma \sum_{s'} p(s'|s, a) v(s'), \forall s \in S, \forall a \in A \end{aligned}$$

The constants α_s are arbitrary positive numbers.

Notes

- Linear programming methods can also be used to solve MDPs
- Linear programming methods become impractical at a much smaller number of states than do DP methods (by a factor of about 100).


Policy iteration: Policy evaluation

- Given a policy π , compute $v_\pi(s)$ for all $s \in S$
- Linear system of $|S|$ equations in $|S|$ unknowns

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_\pi(s')]$$

- Iterative policy evaluation (converge to v_π as $k \rightarrow \infty$)

$$v_{k+1}(s) = \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_k(s')]$$



Policy iteration: Policy Improvement

- Given $v_\pi(s)$, compute $q_\pi(s, a)$

$$q_\pi(s, a) = \sum_{s', r} p(s', r | s, a) [r + \gamma v_\pi(s')]$$

- We can improve π by acting greedily

$$q_\pi(s, \pi'(s)) \geq v_\pi(s)$$

- π' is as good as or better than π

Policy Improvement Theorem

Policy Improvement Theorem

If $\forall s \in \mathcal{S}, q_{\pi}(s, \pi'(s)) \geq v_{\pi}(s)$, then π' is as good as or better than π , i.e., $v_{\pi'}(s) \geq v_{\pi}(s)$.

Greedy policy:

$$\pi'(s) = \arg \max_a q_{\pi}(s, a)$$

Policy iteration: pseudocode

Algorithm 2: Policy Iteration for estimating $\pi \approx \pi_*$

```

1 Initialize  $V(s)$  and  $\pi(s)$  arbitrarily, for all  $s \in \mathcal{S}$ ;
2 while True do
    // Policy Evaluation
3     while  $\Delta \geq \theta$  do
4          $\Delta \leftarrow 0$ ;
5         for each  $s \in \mathcal{S}$  do
6              $v \leftarrow V(s)$ ;
7              $V(s) \leftarrow \sum_a \pi(a|s) \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$ ;
8              $\Delta \leftarrow \max(\Delta, |v - V(s)|)$ ;
    // Policy Improvement
9     policy-stable  $\leftarrow$  True;
10    foreach  $s \in \mathcal{S}$  do
11         $a \leftarrow \pi(s)$ ;
12         $\pi(s) \leftarrow \arg \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$ ;
13        if  $a \neq \pi(s)$  then
14            policy-stable  $\leftarrow$  False;
15    if policy-stable then
16        break;

```


Solving Multi-period inventory problem

We can solve the multi-period inventory problem using policy iteration, value iteration (or Bellman optimality equation, linear programming)

Example:

- State space: $S = \{s \in \mathbb{Z} : -16 \leq s \leq 16\}$
- Action space: $A = \{a \in \mathbb{N} : 0 \leq a \leq 16\}$
- Overage cost: $h = 1$, Underage cost: $p = 10$
- Demand distribution: $D \sim \mathcal{P}(5)$
- Discount factor: $\gamma = 0.9$

Solving Multi-period inventory problem: Example



$$\pi_*(s) = \begin{cases} 0, & 8 \leq s \leq 16 \\ 8 - s, & -8 \leq s < 8 \\ 16, & -16 \leq s < -8 \end{cases}$$

We obtained a **Base-stock policy** for $-8 \leq s \leq 16$, which is exactly the optimal policy for this problem
 ($h = 1, p = 10, D \sim \mathcal{P}(5)$)

Multi-period inventory problem and MDP

Summary

- The multi-period inventory problem is an associative problem
- We need to find the best action for each state
- The multi-period inventory problem can be formulated as an MDP
- Methods for solving MDPs are introduced

Next

- What if we do not know the dynamics of the environment?
- What if the state space/action space is too large to solve?

TD Learning

- Temporal Difference (TD) Learning do not require complete knowledge of the environment.
- Although a model is required, the model need only generate sample transitions.
- In many cases, it is difficult to obtain the distribution in explicit form.
- TD learning:
 - Prediction problem (estimating v_π or q_π)
 - Control problem (estimating π_*)

TD Prediction

- Recall that the value function is the expected return:

$$\begin{aligned}v_{\pi}(s) &= \mathbb{E}_{\pi} [G_t | S_t = s] \\&= \mathbb{E}_{\pi} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]\end{aligned}$$

- The Simplest TD method: TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

- Update $V(S_t)$ towards estimated return $R_{t+1} + \gamma V(S_{t+1})$
- $R_{t+1} + \gamma V(S_{t+1})$ is called TD target
- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called TD error

Q-learning

■ Q-learning

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \left[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right]$$

■ Use ϵ -greedy policy to select action

TD Control (Q-Learning): Pseudocode

Algorithm 3: Q-learning for estimating $\pi \approx \pi_*$

Input: a small $\epsilon > 0$, a small $\alpha \in (0, 1]$

Output: output a deterministic policy $\pi \approx \pi_*$

```
1 Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$ ;  
2 while True do  
3      $t \leftarrow 0$ ;  
4     Initialize  $S_t$ ;  
5     while  $S_t$  is not terminal do  
6         take action  $A_t$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy);  
7         observe  $R_{t+1}$  and  $S_{t+1}$ ;  
8          $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha[R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)]$ ;  
9          $t \leftarrow t + 1$ ;
```

TD Control (Q-Learning): Example

- State space: $S = \{s \in \mathbb{Z} : -16 \leq s \leq 16\}$
- Action space: $A = \{a \in \mathbb{N} : 0 \leq a \leq 16\}$
- Discount factor: $\gamma = 0.9$
- Environment: $h = 1, p = 10, D \sim \mathcal{P}(5)$

	$a_0 = 0$	$a_1 = 1$...	$a_{ A } = 16$
$s_0 = -16$	$Q(s_0, a_0)$	$Q(s_0, a_1)$...	$Q(s_0, a_{ A })$
$s_1 = -15$	$Q(s_1, a_0)$	$Q(s_1, a_1)$...	$Q(s_1, a_{ A })$
...
$s_{ S } = 16$	$Q(s_{ S }, a_0)$	$Q(s_{ S }, a_1)$...	$Q(s_{ S }, a_{ A })$

Table: Q-values for different state-action pairs

TD Control (Q-Learning): Results




Figure: Value iteration




Figure: Q-learning

Multi-period inventory problem and TD learning

Summary

- Q-learning can be used to solve the multi-period inventory problem
- TD learning do not require complete knowledge of the environment
- A model is required to generate sample transitions
- A Q-table is required to store and update Q-values

Next

- The state can be very large or even continuous
- A function approximation method can be used

Deep Q-Network (DQN)

Recall TD error for Q-learning

$$\delta_t = R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a') - Q(S_t, A_t)$$

Basic idea of DQN

- Represent $Q(s, a)$ by a neural network with weights θ

$$Q(s, a; \theta)$$

- Update θ to minimize the TD error

$$\delta_t = R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a'; \theta) - Q(S_t, A_t; \theta)$$

Two issues of DQN

■ Correlations between samples

- Training NN requires independent and identically distributed (i.i.d.) samples
- But, samples taken from an episode are not i.i.d.

■ Non-stationary targets

- The target $y_j = R_{j+1} + \gamma \max_{a'} Q(S_{j+1}, a'; \theta)$ changes as θ changes

$$\delta_t = R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a'; \theta) - Q(S_t, A_t; \theta)$$

Experience replay and target network

■ Experience replay

- Store transitions (S_t, A_t, R_t, S_{t+1}) in a replay buffer D
- Sample random minibatch of transitions (S_j, A_j, R_j, S_{j+1}) from D

■ Target network

- Use a separate network with weights θ^- to compute the target

$$y_j = \begin{cases} r_j & \text{for terminal } S_{j+1} \\ r_j + \gamma \max_{a'} \hat{Q}(S_{j+1}, a'; \theta^-) & \text{for non-terminal } S_{j+1} \end{cases}$$

- Update θ^- every C steps

Deep Q-Network (DQN): Pseudocode

Algorithm 4: Deep Q-Network

Input: replay buffer capacity N , the number of steps C to perform a target update, a small $\epsilon > 0$, a small $\alpha \in (0, 1]$

Output: output a deterministic policy $\pi \approx \pi_*$

```

1 Initialize empty replay memory  $D$  to capacity  $N$ ;
2 Initialize action-value function  $Q$  with random weights  $\theta$ ;
3 Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- \leftarrow \theta$ ;
4 for episode = 1, 2, ...  $M$  do
5      $t \leftarrow 0$ ;
6     Initialize  $S_t$ ;
7     while  $S_t$  is not terminal do
8         with probability  $\epsilon$  select a random action  $A_t$ ;
9         otherwise take action  $A_t$  using policy derived from  $Q$ ;
10        Execute action  $A_t$  and observe reward  $R_t$  and  $S_{t+1}$ ;
11        Store transition  $(S_t, A_t, R_t, S_{t+1})$  in  $D$ ;
12         $t \leftarrow t + 1$ ;
13        Sample random minibatch of transitions  $(S_j, A_j, R_j, S_{j+1})$  from  $D$ ;
14        Set  $y_j = \begin{cases} r_j & \text{for terminal } S_{j+1} \\ r_j + \gamma \max_{a'} \hat{Q}(S_{j+1}, a'; \theta^-) & \text{for non-terminal } S_{j+1} \end{cases}$ ;
15        Perform a gradient descent step on  $(y_j - Q(S_j, a_j; \theta))^2$  with respect to the network parameters  $\theta$ ;
16        Every  $C$  steps reset  $\hat{Q} \leftarrow Q$ ;

```

DRL for inventory management: Example 1




Figure: Perishable Inventory Problem

⁰ De Moor, Bram J., Joren Gijsbrechts, and Robert N. Boute. 2022. "Reward Shaping to Improve the Performance of Deep Reinforcement Learning in Perishable Inventory Management." *European Journal of Operational Research* 301 (2): 535-45.

DRL for inventory management: Summary

Summary

- Machine learning models can be used to approximate the action-value function
- Integrating with machine learning enables solving more complex problems

Future Directions

- Offline Methods
- Explainability
- Privacy-preserving

Table of Contents

- 1 Reinforcement Learning for Inventory Optimization
 - Newsvendor Problem and Multi-armed Bandit Problem
 - Multi-Period Inventory Problem and Markov Decision Process
 - Multi-Period Inventory Problem and TD Learning
 - Inventory Management and Deep Reinforcement Learning
- 2 Surrogate Models in Optimization
 - Surrogate-based Optimization
 - Trust Region Methods
 - Bayesian Optimization
 - Surrogate-assisted Evolutionary Optimization
- 3 Appendix

Surrogate-based Optimization

Motivation:

- The objective function f is highly nonlinear
- The objective function f only can be obtained by simulation

Basic Idea:

- Approximate the objective function f with a surrogate model \hat{f}
- Solve $\mathbf{x} = \arg \min_{\mathbf{x}} \hat{f}(\mathbf{x})$

Note:

- This method only builds surrogate model once
- The accuracy of \hat{f} at the obtained solution is unknown

⁰Gurobi Machine Learning:

Trust Region Methods

Basic Idea:

- Approximate the objective function f with \hat{f} (Normally, a second-order Taylor expansion approximation).
- Solve $\mathbf{x} = \arg \min_{\mathbf{x}} \hat{f}(\mathbf{x})$ within a trust region
- Expand or shrink the trust region based on the improvement
- Accept the new solution if the improvement is large enough

Trust Region Methods: Pseudocode

Algorithm 5: Trust Region Methods

Input: initial guess x_0 , trust region radius δ , threshold η_1, η_2 , scale factor γ_1, γ_2

Output: solution x

```

1  $k \leftarrow 0$ ;
2 solve  $f(x_k)$ ;
3 while not converged do
4   solve  $\min_x \hat{f}(x_k + p)$  s.t.  $\|p\| \leq \delta_k$ ;
5   compute  $\rho_k = \frac{f(x_k) - f(x_k + p_k)}{f(x_k) - \hat{f}(x_k + p_k)}$ ;
6   if  $\rho_k < \eta_1$  then
7      $\delta_{k+1} \leftarrow \gamma_1 \delta_k$ ;
8   else
9      $x_{k+1} \leftarrow x_k + p_k$ ;
10    if  $\rho_k > \eta_2$  then
11       $\delta_{k+1} \leftarrow \gamma_2 \delta_k$ ;
12   $k \leftarrow k + 1$ ;

```

Bayesian Optimization

Motivation:

- Evaluate f is expensive or time-consuming.
- Only a few evaluations of f are allowed.
- No first- or second-order derivatives.

Basic Idea:

- Approximate f with a surrogate model \hat{f} (Gaussian process)
- Find the next evaluation point x by maximizing the acquisition function $\alpha(x)$
- Update \hat{f} with the new evaluation point x and $f(x)$
- Repeat until the budget is exhausted

Gaussian Process

- A set of points $X = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$ is given.
- The corresponding $\mathbf{y} = \{y_1, y_2, \dots, y_n\}$ are observed.
- Predict y_{n+1} at a new point \mathbf{x}_{n+1} .
 - $y_{n+1} | \mathbf{y} \sim \mathcal{N}(\mu(\mathbf{x}_{n+1}), \sigma^2(\mathbf{x}_{n+1}))$
- $\mu(\mathbf{x}_{n+1})$ is the estimated mean of y_{n+1} .
- $\sigma^2(\mathbf{x}_{n+1})$ is the estimated variance of y_{n+1} , which is a measure of uncertainty.

Gaussian Process: Example

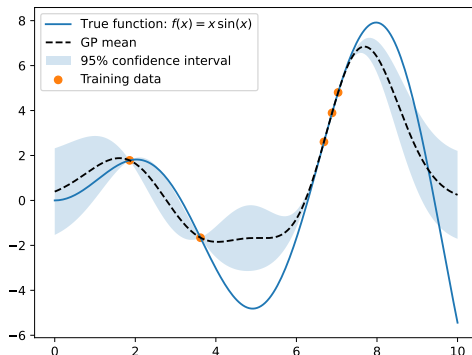


Figure: Gaussian Process Example

Acquisition Function

Find the next evaluation point \mathbf{x}_{t+1} by maximizing $\alpha(\mathbf{x})$.

$$\mathbf{x}_{t+1} = \arg \max_{\mathbf{x}} \alpha(\mathbf{x})$$

- Prediction-based Exploration:

$$\alpha(\mathbf{x}) = -\mu(\mathbf{x})$$


- Error-based Exploration:

$$\alpha(\mathbf{x}) = \sigma(\mathbf{x})$$

- Lower Confidence Bound:

$$\alpha(\mathbf{x}) = -(\mu(\mathbf{x}) - k\sigma(\mathbf{x}))$$

- ...



Bayesian Optimization: Pseudocode

Algorithm 6: Bayesian Optimization

Input: $X = \{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n\}$, $\mathbf{y} = \{y_1, y_2, \dots, y_n\}$, acquisition function $\alpha(x)$, budget T

Output: solution x^*

```
1 while  $t < T$  do
2   Update the Gaussian process model with  $X$  and  $\mathbf{y}$ ;
3    $x_{t+1} \leftarrow \arg \max_x \alpha(x)$ ;
4   Evaluate  $y_{t+1} = f(x_{t+1})$ ;
5   Update  $X = X \cup \{x_{t+1}\}$ ,  $\mathbf{y} = \mathbf{y} \cup \{y_{t+1}\}$ ;
6  $x^* = \arg \min_x \mathbf{y}$ ;
```

Properties of Bayesian Optimization

- $x \in \mathbb{R}^d$, d is not too large. Typically, $d \leq 20$.
- Feasible region \mathcal{X} is a simple set (e.g., a hyper-rectangle).

⁰Frazier, Peter I. "A tutorial on Bayesian optimization." arXiv preprint arXiv:1807.02811 (2018).

Surrogate-assisted Evolutionary Optimization

Motivation:

- The objective function f is expensive . . .
- Evolutionary computation needs many evaluations

Basic Idea:

- Approximate the objective function or constraint functions
 - Regression models
- Predict the feasibility, superiority.
 - Classification models
- Use the surrogate models to evaluate candidate solutions
- Only evaluate the "promising" solutions
- Update the surrogate models with certain frequency

Pseudocode

Algorithm 7: Surrogate-assisted Evolutionary Optimization

Input: $X = \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$, $\mathbf{y} = y_1, y_2, \dots, y_n$, surrogate model M , real function f , budget T

- 1 **while** *not terminate* **do**
 - 2 Update M using X and \mathbf{y} ;
 - 3 Generate new candidate solutions X' ;
 - 4 Evaluate X' with M ;
 - 5 Select promising solutions from X' ;
 - 6 Evaluate the selected solutions with the real function f ;
 - 7 Add the evaluated solutions to X and \mathbf{y} ;
-

Surrogate-assisted Optimization for Inventory

Motivation:

- Difficult to explicitly formulate supply chain simulation models.
- Simulation model under uncertainty requires a large number of samples, which is time-consuming.
- Supply chain digital twin can be expensive to evaluate.




Figure: Supply Chain Digital Twin Developed by anyLogistix

Surrogate-assisted Optimization for Inventory

- Multi-period multi-echelon perishable inventory problem
- Predict total cost/service level of an inventory policy with a surrogate model
 - (s, S) policy, (r, Q) policy, (s, Q) policy, ...
 - e.g., $\mathbf{x} = (s_1, S_1, s_2, S_2, \dots, s_N, S_N)$




Figure: Perishable inventory problem in a distribution system

Surrogate-assisted Optimization for Inventory

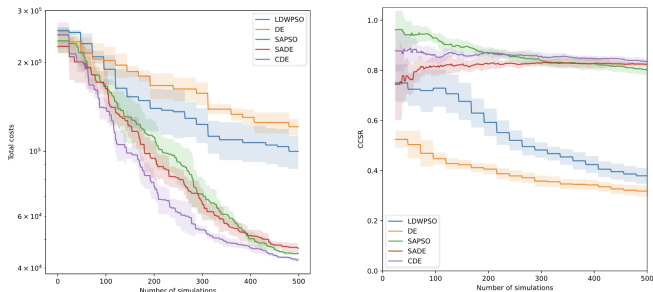


Figure: Convergence graph (Left), correct selection rate (Right)

Summary

Summary:

- Surrogate-assisted optimization is useful when the objective function is expensive.
- Machine learning models can be used to approximate the objective function or constraint functions.
- It is able to obtain a good solution with a small number of evaluations.

Table of Contents

- 1 Reinforcement Learning for Inventory Optimization
 - Newsvendor Problem and Multi-armed Bandit Problem
 - Multi-Period Inventory Problem and Markov Decision Process
 - Multi-Period Inventory Problem and TD Learning
 - Inventory Management and Deep Reinforcement Learning
- 2 Surrogate Models in Optimization
 - Surrogate-based Optimization
 - Trust Region Methods
 - Bayesian Optimization
 - Surrogate-assisted Evolutionary Optimization
- 3 Appendix

Policy Improvement Theorem (Optional)

$$\begin{aligned} v_{\pi}(s) &\leq q_{\pi}(s, \pi'(s)) \\ &= \mathbb{E} [R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s, A_t = \pi'(s)] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma q_{\pi}(S_{t+1}, \pi'(S_{t+1})) | S_t = s] \\ &\leq \mathbb{E}_{\pi'} [R_{t+1} + \gamma R_{t+2} + \gamma^2 q_{\pi}(S_{t+2}, \pi'(S_{t+2})) | S_t = s] \\ &\dots \\ &= v_{\pi'}(s) \end{aligned}$$

Policy iteration: Policy Improvement

- Suppose the new policy π' is as good as, but not better than, the old policy π
- Then, $v_{\pi'}(s) = v_{\pi}(s)$, for all $s \in S$

$$v_{\pi'}(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma v_{\pi'}(s')]$$

- This is the Bellman optimality equation
- π' must be an optimal policy

Value Function ∞ -Norm (Optional)

- We use ∞ -norm to measure the difference between two value functions u and v
- ∞ -norm is the maximum absolute difference between state values

$$\|u - v\|_{\infty} = \max_{s \in S} |u(s) - v(s)|$$

Bellman Expectation Backup is a Contraction (Optional)

- Define the Bellman operator T_π

$$T_\pi(v) = \mathcal{R}_\pi + \gamma \mathcal{P}_\pi v$$

- This operator is γ -Contraction, i.e., it makes values functions closer by at least γ

$$\begin{aligned}\|T_\pi(u) - T_\pi(v)\|_\infty &= \|\mathcal{R}_\pi + \gamma \mathcal{P}_\pi u - \mathcal{R}_\pi - \gamma \mathcal{P}_\pi v\|_\infty \\ &= \gamma \|\mathcal{P}_\pi u - \mathcal{P}_\pi v\|_\infty \\ &\leq \gamma \|u - v\|_\infty\end{aligned}$$

Value iteration

- Must we wait for policy evaluation to converge before improving the policy?
- Policy improvement after one update of each state

$$V_{k+1}(s) = \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V_k(s')]$$

- For arbitrary v_0 , v_k converges to v_*
- The output policy can be obtained

$$\pi(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V_*(s')]$$

Value iteration: pseudocode

Algorithm 8: Value Iteration for estimating $\pi \approx \pi_*$

Input: input a small threshold $\theta > 0$ determining accuracy of estimation

Output: output a deterministic policy $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$$

- 1 Initialize $V(s)$ arbitrarily, for all $s \in \mathcal{S}^+$;
 - 2 Initialize $V(\text{terminal}) = 0$;
 - 3 **while** $\Delta \geq \theta$ **do**
 - 4 $\Delta \leftarrow 0$;
 - 5 **foreach** $s \in \mathcal{S}$ **do**
 - 6 $v \leftarrow V(s)$;
 - 7 $V(s) \leftarrow \max_a \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$;
 - 8 $\Delta \leftarrow \max(\Delta, |v - V(s)|)$;
-

Gaussian Process (Optional)

$$\mu(\mathbf{x}) = \mathbf{K}(\mathbf{x}, X)\mathbf{K}(X, X)^{-1}(\mathbf{y} - m(X)) + m(\mathbf{x})$$

$$\sigma^2(\mathbf{x}_{n+1}) = \mathbf{K}(\mathbf{x}_{n+1}, \mathbf{x}_{n+1}) - \mathbf{K}(\mathbf{x}_{n+1}, X)\mathbf{K}(X, X)^{-1}\mathbf{K}(X, \mathbf{x}_{n+1})$$

Mean function and Kernel:

- $m(\mathbf{x})$: Mean function
- $\mathbf{K}(\mathbf{x}, \mathbf{x}')$: Kernel function

Commonly used functions:

- Zero mean function: $m(\mathbf{x}) = 0$
- Gaussian kernel: $\alpha \exp(-\|\mathbf{x} - \mathbf{x}'\|^2)$
- ...